

What an Agent Workspace Reveals About the Abstractions Its Operating System Is Missing

Anonymous Author(s)

ABSTRACT

AI agents are becoming long-lived services that plan, invoke tools, and persist state, yet they are still hosted by application-layer workspaces built on operating-system abstractions drawn for static workloads. We built and ran such a workspace—an open-source, multi-agent chat application with persistent per-agent memory—and report where its implementation hit walls that are not application problems but missing-primitive problems: at each wall the workspace was forced into a degraded workaround. We enumerate six workarounds, measure two, and map each onto an OS abstraction that would remove it. The contribution is not a new OS; it is consumer-side evidence for the set of abstractions the agent era is starved for.

ACM Reference Format:

Anonymous Author(s). 2026. What an Agent Workspace Reveals About the Abstractions Its Operating System Is Missing. In *Proceedings of ACM Conference (Conference'17)*. ACM, New York, NY, USA, 2 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 INTRODUCTION

Operating systems expose processes, threads, files, sockets, and resource controllers as the units of execution, composition, and accounting. These abstractions were drawn for workloads that are static, predictable, and short-lived relative to human attention. AI agents invert all three properties: they are long-lived, semantically adaptive, and interact with tools and environments in ways the runtime cannot predict. A recent workshop series argues that this mismatch is not incidental but structural, and that operating systems themselves must become *agentic* [1].

Most of the evidence offered for this claim comes from the *producer* side of the OS interface: new kernel primitives, new schedulers, new isolation models. This paper offers evidence from the *consumer* side. We built and operate a real, open-source agent workspace—a multi-agent chat application in which a human routes work to named agents via channel mentions, and each agent retains memory across sessions. The application is small (a few hundred lines) and states its non-goals explicitly. We found that at every hard problem, the limiting factor was not the application's logic but the absence of an operating-system abstraction that should plausibly exist.

The contribution is the mapping. Each of the workspace's admitted workarounds corresponds to an OS-level abstraction that

would dissolve it. We measure two of these workarounds to show they are not cosmetic: a serial execution queue whose latency grows linearly with concurrency, and a prompt-replay memory scheme whose input-token cost grows linearly with history length. Each is technically re-implementable in user space, but only by re-building, per application, the scheduling and managed-state primitives that belong one layer down—and by getting their corner cases right every time.

We are explicit about what we are *not* claiming: we do not propose a new operating system, nor argue that any particular abstraction is the correct one. We argue that a real application, by being honest about its walls, enumerates the abstractions the agent stack lacks.

2 THE WORKSPACE

The system is an open-source agent workspace (anonymized for review). A human types into a single chat channel; messages beginning with @Name are routed to the named agent. Each agent has a persona and a private, append-only memory. On each turn the agent's full history is replayed into the prompt, the model is invoked through a standard OpenAI-compatible endpoint, and the reply is stored back. Memory is isolated per agent: a fact told to one agent is not visible to another.

The implementation comprises a Node.js HTTP router, a SQLite store for channel history and per-agent memory, and a minimal web frontend. It uses the host Node runtime's built-in SQLite module; no external agent framework, vector database, or daemon is involved. The system is deployed and used.

What makes the system useful as evidence is not its capability but the explicitness of its non-goals. By design, it has:

- **No multi-agent concurrency.** Messages in a channel are processed through one serial queue. Two agents cannot act on the same state at the same time.
- **No memory compression.** The entire history is replayed into the prompt on every turn. There is no summarization or archival layer.
- **No inter-agent isolation.** All agents share one process and one database; there is no boundary preventing one agent from reading another's state.
- **No capability routing.** A mention is parsed by a regular expression; there is no matching of work to the agent best suited for it, and no delegation between agents.

These are not oversights. Each is a deliberate non-goal that kept the implementation readable. We now argue that each is also a place where an operating-system abstraction is conspicuously absent.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference'17, July 2017, Washington, DC, USA

© 2026 Association for Computing Machinery.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

Table 1: Each workaround maps to an OS abstraction that would remove it.

Workaround (what is hand-rolled)	OS abstraction wanted
Serial queue (no concurrency)	Semantics-aware agent scheduling
Full-history replay (no compression)	Long-lived managed state
One process + one DB (no boundary)	Execute-only / attested isolation
JSON personas (hand-authored)	Skills as unit of composition
@Name regex (no matching)	Agent-facing interface / IPC
SQLite file (no managed object)	Agent state as OS object

3 WORKAROUNDS AND THE ABSTRACTIONS THEY WANT

Table 1 lists the six workarounds and the OS abstraction each implies. Two are measured below.

3.1 Measured: the serial queue

Because the workspace provides no agent scheduling, concurrent requests to one channel serialize. We fired N concurrent mentions at a channel for $N \in \{1, 2, 4, 8\}$ and measured end-to-end latency. Wall-clock time grew linearly: 3.3 s, 7.0 s, 13.9 s, 27.3 s for $N = 1, 2, 4, 8$ —the eighth request waited over 27 seconds for seven predecessors. This is not a slow model; it is a *degenerate scheduler* that imposes a global order the application has no basis to choose. An OS that accounted for agents as first-class schedulable entities—with their intent, priority, and exploration semantics—could reorder, batch, or interleave these requests. The application can only fall back to a queue because no such abstraction exists beneath it.

3.2 Measured: prompt-replay memory

Because the workspace provides no managed memory abstraction, it replays the full conversation history into the prompt on every turn. We held an 8-turn conversation with one agent and read the model’s reported input token count per turn. Prompt tokens grew 39, 81, 118, 142, 173, 206, 240, 308—tracking history length linearly, so each turn re-pays for every prior turn. An agent that persists across days or weeks, as the AgenticOS position assumes, becomes costly under this scheme well before it becomes unusable. The fix is not a better application prompt; it is a memory object the runtime manages: episodic storage with compression, recall, and eviction policies that the application requests rather than implements.

4 WHY THESE ARE ABSTRACTIONS, NOT FEATURES

A reviewer may object that the items in Table 1 are missing *features*, not missing *abstractions*. The distinction is the contribution. A feature is private to one application; an abstraction is a contract many applications share and that the layer beneath them guarantees.

Each row passes this test. Any agent workspace—not only ours—that lets agents persist, collaborate, or be selected by capability will hit the same wall, because the wall is in the contract between the

application and its runtime, not in the application’s logic. Concurrency between agents is a scheduling contract; durable, compressible context is a managed-state contract; preventing one agent from reading another is an isolation contract; routing work by capability is an interface contract. These are the shapes of OS abstractions. The fact that a few-hundred-line application is forced to hand-roll all of them, badly, is the evidence.

5 RELATION TO THE AGENTICOS AGENDA

The mapping aligns with topics the AgenticOS community has already identified from the producer side. Agent-aware resource control [2] corresponds to our scheduling row; long-lived state abstractions for agent context and memory [1] correspond to our replay row; execute-only and attested-channel isolation [4, 5] corresponds to our isolation row; skills as a composition unit [3] corresponds to our persona row; agent-facing interface redesign [6] corresponds to our routing row. The agreement is the point: the abstractions an application is starved for are the same ones the OS community is beginning to build. Consumer-side evidence corroborates the producer-side agenda.

6 WHAT WE ARE NOT CLAIMING

We do not claim a new operating system. We do not claim that the abstractions named in Table 1 are the *correct* ones, only that they are *wanted*. We do not claim the workspace is representative of all agent applications; we claim that any application with its shape—persistent, multi-agent, capability-routed—will hit the same walls because the walls are in the runtime contract. Finally, we do not claim these abstractions are easy; we claim their absence is expensive, and we have the measurements to show it.

7 CONCLUSION

An application-layer agent workspace that states its non-goals explicitly is a measuring instrument for the abstractions its operating system lacks. We built one, measured two of its walls, and mapped six to OS-level abstractions the agent era needs. The hope is that consumer-side evidence of this kind complements the producer-side primitives the community is beginning to propose, and that the contract between agent applications and their runtime becomes something the application requests rather than re-implements.

REFERENCES

- [1] AgenticOS Workshop Series. *Operating Systems Design for AI Agents*. SOSP 2026 (2nd edition); ASPLOS 2026 (1st edition). <https://os-for-agent.github.io/>
- [2] Y. Zheng, J. Fan, Q. Fu, Y. Yang, W. Zhang, A. Quinn. *AgentCgroup: Understanding and Controlling OS Resources of AI Agents*. AgenticOS @ ASPLOS 2026.
- [3] L. Chen, Z. Wang, W. Zheng, E. Feng, D. Du, Y. Xia, H. Chen. *Skills are the new Apps—Now It Is Time for Skill OS*. AgenticOS @ ASPLOS 2026.
- [4] R. Tiwari, D. Williams. *Execute-Only Agents: Architectural Defense Against Prompt Injection for AI Agents*. AgenticOS @ ASPLOS 2026.
- [5] Q. Wu, W. Zhang, G. Fang, S. Mao, B. Gao, D. Levitsky, S. M. Butterworth, R. Cameron. *Grimlock: Guarding High-Agency Systems with eBPF and Attested Channels*. AgenticOS @ ASPLOS 2026.
- [6] Y. Wang, M. Li, H. Chen. *Rethinking OS Interfaces for LLM Agents*. AgenticOS @ ASPLOS 2026.